
Inductive Bias Meta-Learning with Generative Models

Name Surname
MIPT
Moscow, Russia
email@phystech.edu

Name Surname
MIPT
Moscow, Russia
email@phystech.edu

Abstract

This paper investigates inductive bias in machine learning models. By inductive bias we mean the preference of a model for certain types of functions or data structures over others. To analyze the inductive bias of a fixed model, we consider a problem of finding data that this model can fit and generalize on particularly well. Previous work demonstrated that generating labels for a fixed dataset allows one to extract the inductive bias. Here, we extend this approach by proposing a method for generating full synthetic datasets. We train a generative model to produce datasets on which the target model achieves strong generalization performance. We test the proposed framework on CNNs and RNNs, and analyze obtained datasets for each model.

Keywords: inductive bias, meta-learning, generative models.

1 Introduction

Inductive bias plays a fundamental role in a model’s ability to generalize [1]. Without additional assumptions, any hypothesis that fits a finite training dataset is considered equally likely by the model, which can lead to poor generalization. Thus, effective generalization requires a set of prior assumptions about the data, and these assumptions constitute the model’s inductive bias. Understanding and exploiting inductive bias can be a powerful tool for enhancing model performance.

The term “inductive bias” does not have a single, strict definition. For instance, inductive bias can mean underlying assumptions about the domain and task or specific implementation choices that constrain or order the space of hypotheses [2]. In this paper, by inductive bias we mean a model’s preference for certain types of functions or data structures over others, expressed through better generalization performance. For example, convolutional neural networks (CNNs) are inductively biased towards data with local structure and translation invariance [3–5], while recurrent neural networks (RNNs) are inductively biased towards ordered data with sequential structure [6–8].

A clear understanding of the correspondence between models (or their specific components) and inductive biases is useful for solving machine learning problems. However, this task is highly non-trivial: for most models with complex architecture, the problem of detecting inductive bias is analytically intractable. Several approaches have been developed to extract and characterize the inductive bias of models. Boopathy et al. [9] proposed a method to quantify the amount of inductive bias of tasks or models as the amount of information required to specify well-generalizing models within a given hypothesis space. Si et al. [10] investigated the inductive bias of in-context learning by constructing underspecified demonstrations. Li et al. [11] proposed a framework to capture the inductive biases in a learning system by meta-learning Gaussian process kernel hyperparameters from its predictions.

One recent result in this area is the meta-learning framework introduced by Dorrell et al. [12], which identifies the inductive bias of differentiable supervised learning algorithms by generating labels for a

fixed dataset. For all its strengths, this approach has certain limitations. First, it requires access to input data or its distribution. Second, its analysis of inductive bias is constrained by the fixed dataset and its structure.

Here, we propose to extend this approach by generating entire synthetic datasets. The meta-learning loop proceeds as follows: an outer generative model, or meta-learner, produces a dataset, which is used to train and test an inner model, or learner, whose inductive bias we aim to extract. After the learner is trained, the quality of its predictions on test data from this dataset is incorporated into the meta-learner’s loss function. Thus, the meta-learner learns to generate datasets on which the target model shows the best generalization performance; these are datasets towards which it is inductively biased.

The proposed approach requires no knowledge of the input data distribution, making it more flexible. We expect it to enable a more comprehensive analysis of inductive bias.

Contributions. Our contributions can be summarized as follows:

- We present a generative meta-learning framework for extracting the inductive bias of a target model. Unlike prior work [12] that generates only labels for samples from a fixed distribution, our approach produces entire synthetic datasets via a meta-learned generative model.
- We empirically evaluate our algorithm on CNNs and RNNs and compare the results with the known inductive biases of these architectures. For CNNs, we expect to detect a local structure in the generated datasets; for RNNs, we expect them to have a sequential structure.
- We highlight the implications of our findings for understanding and leveraging inductive bias. Our method is general-purpose: it can be used to extract inductive bias of any differentiable supervised learning algorithm and, consequently, to improve the generalization performance of models in machine learning problems.

Outline. The rest of the paper is organized as follows...

2 Problem statement

We study the problem of extracting and characterizing the inductive bias of learning models by finding the datasets that they generalize best. We now present a formal statement of the problem.

Let $\mathbb{X} \subset \mathbb{R}^n$ be the input space, $\mathbb{Y} \subset \mathbb{R}^m$ be the output space and $\mathbf{f}_\theta : \mathbb{X} \rightarrow \mathbb{Y}$ be the target model, where $\theta \in \Theta \subset \mathbb{R}^k$ are its parameters. Let $\ell(\mathbf{f}_\theta(\mathbf{x}), \mathbf{y})$ be the loss function of the target model.

Consider a parametric family of distributions $p(\mathbf{x}|\mathbf{y}, \mathbf{w})$, where $\mathbf{w} \in \mathbb{W} \subset \mathbb{R}^l$ are the parameters of the distribution. For a fixed label $\mathbf{y} \in \mathbb{Y}$, a sample $\mathbf{x} \sim p(\mathbf{x}|\mathbf{y}, \mathbf{w})$ is produced by the generator $\mathbf{g}(\mathbf{w})$.

We aim to find data distribution function p^* , corresponding to parameters $\mathbf{w}^*$, the solution to the following optimization problem:

$$\mathbf{w}^* = \arg \min_{\mathbf{w} \in \mathbb{W}} \mathbb{E}_{\mathcal{D}_{\text{train}}, \mathcal{D}_{\text{test}} \sim p(\mathbf{x}|\mathbf{y}, \mathbf{w})} [\mathcal{L}(\mathbf{f}_\theta, \mathcal{D}_{\text{test}})] \quad \text{s.t.} \quad |\mathcal{D}_{\text{train}} \cup \mathcal{D}_{\text{test}}| = d$$

where

$$\hat{\theta} = \arg \min_{\theta \in \Theta} \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}_{\text{train}}} \ell(\mathbf{f}_\theta(\mathbf{x}), \mathbf{y}).$$

Function $\mathcal{L}$ measures a model’s ability to generalize by calculating the prediction error on a test set and also contains regularization terms to avoid trivial solutions. For example, prediction error can be estimated as cross-entropy or mean squared error (MSE), defined respectively as:

$$\mathcal{L}_{\text{CE}} = -\frac{1}{|\mathcal{D}_{\text{test}}|} \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}_{\text{test}}} \sum_{i=1}^m y_i \log([\mathbf{f}_{\hat{\theta}}(\mathbf{x})]_i), \quad \mathcal{L}_{\text{MSE}} = \frac{1}{|\mathcal{D}_{\text{test}}|} \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{D}_{\text{test}}} \|\mathbf{y} - \mathbf{f}_{\hat{\theta}}(\mathbf{x})\|^2.$$

The regularization term can be the entropy of the found distribution, defined as:

$$\mathcal{R}_{\text{ent}}(\mathbf{w}) = -\mathbb{E}_{\mathbf{y}} \left[\int_{\mathbb{X}} p(\mathbf{x}|\mathbf{y}, \mathbf{w}) \log p(\mathbf{x}|\mathbf{y}, \mathbf{w}) d\mathbf{x} \right].$$

3 Method

Our proposed approach is to meta-train a generative model to produce datasets on which the target model exhibits the best generalization performance, i.e., datasets to which it is inductively biased. In this section we describe our method in detail.




Figure 1: Meta-learning loop

The meta-learning cycle is structured as follows. First, the generator $g(\mathbf{w})$ with current parameters $\mathbf{w}$ constructs a synthetic dataset $\mathcal{D}_{\text{syn}}$ according to the predefined configuration: for each class, a fixed number of samples is generated. The number of classes and the number of samples per class are both specified as hyperparameters. Then, the resulting dataset is split into training $\mathcal{D}_{\text{train}}$ and test $\mathcal{D}_{\text{test}}$ sets.

Next, the target model f_{θ} is trained on the training set $\mathcal{D}_{\text{train}}$ in the inner loop. It is important that the model be re-initialized from scratch at each meta-learning iteration. At the output of the inner loop we obtain a trained model $f_{\hat{\theta}}$.

Next, the meta-loss function is calculated. It consists of two terms. The first term $\mathcal{L}_g$ measures the model’s generalization quality, estimated on the test set $\mathcal{D}_{\text{test}}$. This could be, for example, MSE or Cross-Entropy. The second term $\mathcal{R}$ is the regularization term. Note that if all objects within a class are identical, the target model would easily

achieve 100% accuracy. However, this degenerate case is uninformative. The regularization term is precisely what prevents such trivial solutions. Possible choices include the average distance between objects within the same class, defined as:

$$\mathcal{R}_{\text{dist}} = \frac{1}{C} \sum_{k=1}^C \frac{1}{n_k(n_k - 1)} \sum_{\substack{i,j \in \mathcal{I}_k \\ i \neq j}} \|\mathbf{x}_i - \mathbf{x}_j\|_2^2,$$

where C is the number of classes in the dataset, n_k is the number of objects in class k , and $\mathcal{I}_k$ is the set of indices of objects belonging to class k .

Another regularization strategy is to penalize small singular values of the generator’s weight matrix, as this helps to avoid degeneracy. The corresponding regularization term is defined as:

$$\mathcal{R}_{\text{svd}} = \frac{1}{r} \sum_{k=1}^r (\max(\tau - \sigma_k, 0))^2,$$

where $\sigma_1, \dots, \sigma_r$ are the singular values of the weight matrix $\mathbf{X} \in \mathbb{R}^{n \times m}$, $r = \min(n, m)$, and τ is a fixed threshold.

Thus, during meta-training, the meta-loss is minimized by updating the generator parameters $\mathbf{w}$, resulting in a generator $g(\hat{\mathbf{w}})$ that produces datasets which are easy for the target model to generalize from. The structure of these datasets reflects the model’s inductive bias.

The proposed meta-learning procedure is summarized in Algorithm 1.

Algorithm 1 Meta-learning for inductive bias extraction

- 1: Initialize generator $g(\mathbf{w})$ with random parameters $\mathbf{w}$
 - 2: **while** step < total steps **do**
 - 3: Generate synthetic dataset $\mathcal{D}_{\text{syn}} = \{(\mathbf{x}_i, \mathbf{y}_i)\}$ with C classes and N samples per class
 - 4: Split $\mathcal{D}_{\text{syn}}$ into training set $\mathcal{D}_{\text{train}}$ and test set $\mathcal{D}_{\text{test}}$
 - 5: Re-initialize target model f_{θ} from scratch
 - 6: Train f_{θ} on $\mathcal{D}_{\text{train}}$ in the inner loop $\rightarrow$ trained model $f_{\hat{\theta}}$
 - 7: Compute meta-loss: $\mathcal{L}_{\text{meta}} = \mathcal{L}_g(f_{\hat{\theta}}, \mathcal{D}_{\text{test}}) + \lambda \mathcal{R}(\mathcal{D}_{\text{syn}})$
 - 8: Take $\mathbf{w}$ gradient step on meta-loss
 - 9: **end while**
-

Similar to previous work [12], our approach can be extended to compare the inductive biases of two models, i.e., to find a dataset that is easy for one model to generalize from but difficult for the other. In this case, both models are trained in the inner loop, and the measure of the generalization quality of the second model is subtracted when computing the meta-loss. Consequently, minimizing this meta-loss improves the generalization of the first model while degrading that of the second.

4 Experiments

4.1 Baseline experiment

First, we replicate the experiment by Dorell et al. [12] on extracting inductive bias via label generation. Their algorithm proceeds as follows: the meta-learner generates labels for a dataset, which is subsequently used to train the learner; the learner’s generalization performance on these labels is then evaluated and incorporated into the meta-loss function. In this way, the meta-learner is trained to find functions that the learner can easily generalize.

As the target model, we use linear regression, and as the meta-learner, a 4-layer MLP. The meta-loss function consists of three terms: the MSE of the linear regression on the test set, the Sinkhorn distance of the label distribution from a uniform distribution, and terms that account for previously obtained functions to ensure orthogonality with them.

As a result of the experiment, we obtained three orthogonal linear functions, as expected.

4.2 Proposed Framework Evaluation

To empirically validate our algorithm, we conducted computational experiments to extract the inductive bias of CNNs and RNNs. Specifically, we verify that the generative model learns to create datasets on which the target model shows the best generalization performance, investigate the influence of different regularization terms on the generated data, and attempt to interpret the obtained results in light of the well-known inductive biases of these architectures.

Data in our experiments is fully synthetic. For CNNs we generate datasets of images with a given number of classes and number of objects per class. As a generative model, we used a variational autoencoder (VAE) decoder. It transforms a vector from a latent space, consisting of class embedding and random noise sampled from a standard normal distribution $\mathcal{N}(0, 1)$, into a grayscale image.

The meta-learning process is implemented using the learn2learn library [13]. Specifically, to enable end-to-end differentiability, we wrap the inner learner in a MAML (Model-Agnostic Meta-Learning) [14] wrapper. This allows us to clone the learner for each meta-step, train it on the generated dataset, and still keep the computational graph needed to propagate gradients from the meta-loss back to the generator. We use Adam [15] as the optimizer for meta-learning.

As measures of generalization performance, we use MSE and cross-entropy and then compare them. We also conduct a comparative analysis of different types of regularization terms: the entropy of the distribution, the average distance between objects of the same class, and the singular values of the generator’s weight matrix.

List of expected figures:

- graph comparing different types of regulation (x-axis - meta-iterations, y-axis - accuracy, %);
- figure with examples of the resulting datasets;
- meta-loss change chart (x-axis - meta-iterations, y-axis - meta-loss).

4.3 CNNs and 1×4 Images

We first test our method on a simple example: CNNs with a 1×2 convolution and 1×4 images. As a generator, we use a trainable tensor, meaning that the elements of the synthetic dataset are directly optimized as parameters. For the meta-learning process, we wrap the CNN in a MAML (Model-Agnostic Meta-Learning) wrapper [14] from the learn2learn library [13]. For the outer loop, we use the Adam optimizer [15].

The CNN architecture is as follows: a convolution with a 1×2 kernel and stride (1,2), followed by max pooling, and finally a linear layer that outputs class logits. In the experiment we use the following parameters: 2 classes, 20 samples per class, 5 inner loop steps, and learning rates of 0.5 and 0.01 for inner and outer loop respectively. As a measure of prediction error on the test set, we use MSE. For regularization, we experiment with two terms: the mean intra-class distance and a penalty on small singular values of the data matrix computed via SVD.



Figure 2: Obtained dataset without regularization



Figure 3: Obtained dataset with class diversity regularization



Figure 4: Obtained dataset with SVD regularization

As can be seen, even without regularization, the trivial solution — where all images within a class are identical — does not fully occur. This reveals the convolutional architecture’s inductive bias towards locality: within a given class, we observe images where the first and second pairs of pixels are swapped. This indicates that for this architecture, local structure matters more than global structure.

Adding the mean intra-class distance as a regularization term makes the classes more diverse while preserving locality. In class 1, the defining pattern is two black pixels in one of the two positions, while the other position can contain anything — two white pixels, one white and one black, or vice versa. Moreover, the two black pixels may appear either in the left or the right pair. This reflects the inductive bias towards locality: the label is determined by the local structure rather than the global one. The effect of stride (1,2) is also noticeable — class 0 contains an element with two black pixels in the middle, but this does not cause it to be misclassified as class 1.

When penalizing small singular values, the situation is similar — here, the distinguishing pattern for class 0 consists of a black pixel followed by a white pixel, in that exact order.

Thus, the structure of the obtained datasets reflects the inductive bias of CNNs toward data with local structure, which is consistent with previously known results.

4.4 CNNs and 4x4 images

Next, we conduct an experiment with CNNs and 4×4 images, exploring different combinations for the meta-loss function. For prediction error, we use MSE and cross-entropy; for regularization, we use mean intra-class distance and a penalty on small singular values of the data matrix.

In this experiment, we use a variational autoencoder (VAE) decoder as the generative model. It transforms a vector from a latent space — consisting of a class embedding and random noise sampled from a standard normal distribution $\mathcal{N}(0, 1)$ — into a grayscale 4×4 image.

The CNN consists of two convolutional blocks, each followed by ReLU activation, with a max pooling layer after the first block. The extracted features are then flattened and passed through two fully connected layers to produce class logits.

The experimental parameters are as follows: latent space dimension of 32, 5 classes, 20 samples per class, 100 inner loop steps, and learning rates of 0.1 and 0.01 for the inner and outer loops respectively.

As expected, without regularization terms, we observe the trivial case: all images within each class are identical. The results for the remaining configurations are presented in figures 6–9.




Figure 5: Comparison of different regularization strategies




Figure 6: Obtained dataset for Cross-entropy with class diversity regularization




Figure 7: Obtained dataset for Cross-entropy with SVD regularization




Figure 8: Obtained dataset for MSE with class diversity regularization




Figure 9: Obtained dataset for MSE with SVD regularization

Figure 10 shows a comparative analysis of accuracy as a function of meta-iteration for different meta-loss configurations. The fastest convergence to 100% accuracy occurs without regularization. With mean intra-class distance regularization, convergence is nearly as fast, though the solution is slightly less stable. With SVD regularization, convergence is slower, but the resulting solution is more stable.

To characterize the properties of the obtained datasets, additional analysis is required. TODO




Figure 10: Comparison of different meta-loss strategies

5 Discussion

6 Conclusion

References

- [1] Tom M Mitchell. The need for biases in learning generalizations. 1980.
- [2] Foster John Provost and Bruce G Buchanan. Inductive policy: The pragmatics of bias selection. *Machine Learning*, 20(1):35–61, 1995.
- [3] Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998. doi: 10.1109/5.726791.
- [4] Zihao Wang and Lei Wu. Theoretical analysis of inductive biases in deep convolutional networks, 2024. URL <https://arxiv.org/abs/2305.08404>.
- [5] Nadav Cohen and Amnon Shashua. Inductive bias of deep convolutional networks through pooling geometry, 2017. URL <https://arxiv.org/abs/1605.06743>.
- [6] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural Comput.*, 9(8):1735–1780, November 1997. ISSN 0899-7667. doi: 10.1162/neco.1997.9.8.1735. URL <https://doi.org/10.1162/neco.1997.9.8.1735>.
- [7] Taiga Ishii, Ryo Ueda, and Yusuke Miyao. Empirical analysis of the inductive bias of recurrent neural networks by discrete fourier transform of output sequences, 2023. URL <https://arxiv.org/abs/2305.09178>.
- [8] Igor Dubinin and Felix Effenberger. Fading memory as inductive bias in residual recurrent networks. *Neural Networks*, 173:106179, 2024. ISSN 0893-6080. doi: <https://doi.org/10.1016/j.neunet.2024.106179>. URL <https://www.sciencedirect.com/science/article/pii/S0893608024001035>.
- [9] Akhilan Boopathy, William Yue, Jaedong Hwang, Abhiram Iyer, and Ila Fiete. Towards exact computation of inductive bias, 2024. URL <https://arxiv.org/abs/2406.15941>.
- [10] Chenglei Si, Dan Friedman, Nitish Joshi, Shi Feng, Danqi Chen, and He He. Measuring inductive biases of in-context learning with underspecified demonstrations, 2023. URL <https://arxiv.org/abs/2305.13299>.

- [11] Michael Y Li, Erin Grant, and Thomas L Griffiths. Meta-learning inductive biases of learning systems with gaussian processes. In *Fifth Workshop on Meta-Learning at the Conference on Neural Information Processing Systems*, 2021.
- [12] William Dorrell, Maria Yuffa, and Peter Latham. Meta-learning the inductive biases of simple neural circuits, 2023. URL <https://arxiv.org/abs/2211.13544>.
- [13] Sébastien M. R. Arnold, Praateek Mahajan, Debajyoti Datta, Ian Bunner, and Konstantinos Saitas Zarkias. learn2learn: A library for meta-learning research, 2020. URL <https://arxiv.org/abs/2008.12284>.
- [14] Chelsea Finn, Pieter Abbeel, and Sergey Levine. Model-agnostic meta-learning for fast adaptation of deep networks, 2017. URL <https://arxiv.org/abs/1703.03400>.
- [15] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization, 2017. URL <https://arxiv.org/abs/1412.6980>.

A Appendix / supplemental material

A.1 Additional experiments / Proofs of Theorems

TODO